

# Machine Learning and Data Mining (IT4242E)

**Quang Nhat NGUYEN**

*quang.nguyennhat@hust.edu.vn*

---

Hanoi University of Science and Technology  
School of Information and Communication Technology  
Academic year 2021-2022

# The course's content:

- Introduction
- **Performance evaluation of the ML/DM system**
- Regression problem
- Classification problem
- Clustering problem
- Association rule mining problem

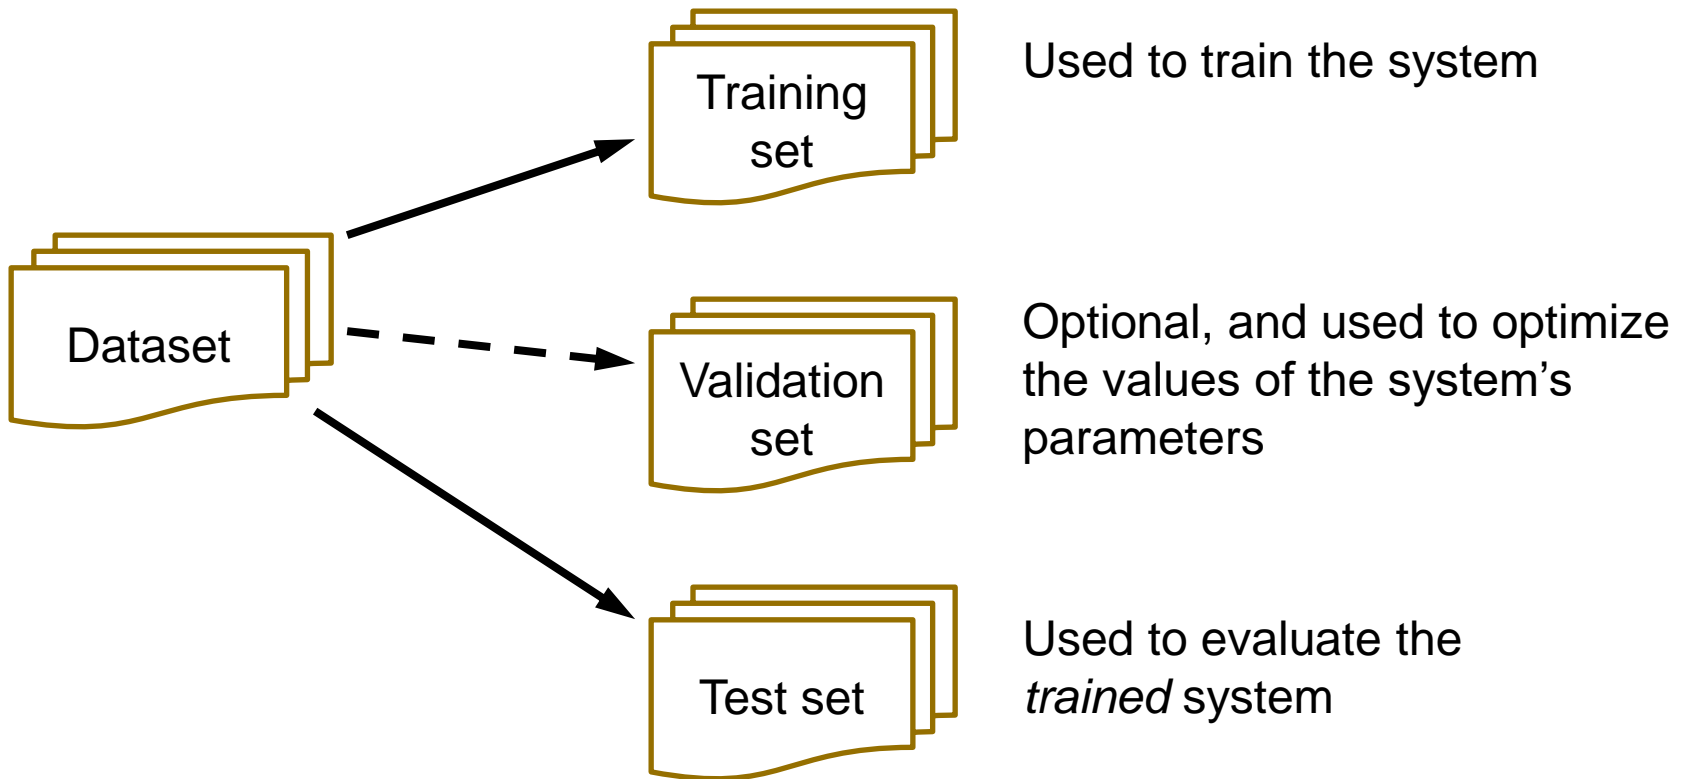
# Performance evaluation (1)

- The evaluation of the performance of a ML or DM system is usually done **experimentally** rather than analytically
  - An analytical evaluation aims at proving a system is correct and complete (e.g., theorem provers in Logics)
  - *But, it is impossible to build a formal definition of a problem to be solved by a ML or DM system (For a ML or DM problem, what are correctness and completeness?)*

# Performance evaluation (2)

- The evaluation of the system performance should:
  - *Be done automatically by the system*, by using a set of test examples (i.e., a test set)
  - Not involve any test users
- Evaluation **methods**
  - How to have a convincing/confident evaluation of the system performance?
- Evaluation **metrics**
  - How to measure (i.e., to compute) the performance of the system?
  - Different metrics for different types of problems (e.g., classification, regression, clustering)

# Evaluation methods (1)



# Evaluation methods (2)

- How to get a confident/convincing evaluation of the system performance?
  - The larger the training set is, the higher the performance of the trained system is
  - The larger the test set is, the more confident/convincing the evaluation is
  - Problem: Very difficult (i.e., rarely) to have (very) large dataset(s)
- *The system performance depends on not only ML/DM algorithms used, but also:*
  - Class distribution
  - Cost of misclassification
  - Size of the training set
  - Size of the test set

# Evaluation methods (3)

- Hold-out (Splitting)
- Stratified sampling
- Repeated hold-out
- Cross-validation
  - $k$ -fold
  - Leave-one-out
- Bootstrap sampling

# Hold-out (Splitting)

- The whole dataset  $D$  is divided into 2 **disjoint subsets**
  - Training set  $D_{train}$  – To train the system
  - Test set  $D_{test}$  – To evaluate the performance of the trained sys.  
→  $D = D_{train} \cup D_{test}$ , and usually  $|D_{train}| \gg |D_{test}|$
- Requirements:
  - Any examples in the test set  $D_{test}$  must not be used in the training of the system
  - Any examples used in the training of the system (i.e., those in  $D_{train}$ ) must not be used in the evaluation of the trained system
  - The test examples in  $D_{test}$  should allow an unbiased evaluation of the system performance
- Usual splitting:  $|D_{train}| = (2/3) \cdot |D|$ ,  $|D_{test}| = (1/3) \cdot |D|$
- **Suitable if we have a large dataset  $D$**



# Stratified sampling

- For such datasets that is small (in size) or unbalanced, the examples in the training and test sets may not be representative
- For example: There are (very) few examples for a specific class label
- Goal: The class distribution in the training and test sets should be approximately equal to that in the original dataset ( $D$ )
- Stratified sampling
  - An approach to have a balanced (in class distribution) dataset
  - Guarantee the class distributions (i.e., the percentages of examples for class labels) in the training and tests set are approximately equal
- The stratified sampling method can not be applied to a regression problem (because for that problem the system's output is a real value, not a discrete value / class label)

# Repeated hold-out

- To apply the Hold-out evaluation method for multi times (i.e., multi runs), each one uses a different training and test sets
  - For each run, a certain percentage of the dataset  $D$  ***is randomly selected*** to create the training set (possibly together with the stratified sampling method)
  - The error values (or the values of other measure metrics) *are averaged* amongst the runs to get the final (average) error value
- This evaluation method is still not perfect
  - In each run, a different test set is used
  - There are still some overlapping (i.e., repeatedly) used examples among those test sets

# Cross-validation

- To avoid any overlapping amongst the used test sets (i.e., the same examples are contained in some different test sets)
- $k$ -fold cross-validation
  - The whole dataset  $D$  is divided into  $k$  **disjoint subsets** (called “*fold*”) that have approximately equal sizes
  - For each run (i.e., of the total  $k$  runs), a subset is circulated to use for the test set, and the remaining  $(k-1)$  subsets are used for the training set
  - The  $k$  error values (i.e., each one for each *fold*) are averaged to get the overall error value
- Usual choices of  $k$ : 10, or 5
- Often, each subset (i.e., fold) is stratified sampling (i.e., to approximate the class distribution) prior to apply the Cross-validation evaluation method
- Suitable if we have a small to medium dataset  $D$

# Leave-one-out cross-validation

- A type of the Cross-validation method
  - The number of folds is exactly the size of the original dataset ( $k=|D|$ )
  - Each fold contains just one example
- To maximally exploit the original dataset
- No random sub-sampling
- Not possible to apply the stratified sampling method
  - Because in each run (loop), the test set contains just one example
- (Very) high computational cost
- Suitable if we have a (very) small dataset  $D$

# Bootstrap sampling (1)

- The Cross-validation method applies sampling without replacement  
→ For each example, *once selected (used) for the training set, then it cannot be selected (used) again (one more time) for the training set*
- The Bootstrap sampling method applies **sampling with replacement** to create the training set
  - Assume that the whole dataset  $D$  contains  $n$  examples
  - To sample with replacement (i.e., repeating) for  $n$  times for the dataset  $D$  to create the training set  $D_{train}$  that contains  $n$  examples
    - From the dataset  $D$ , randomly select an example  $x$  (but **not remove**  $x$  from the dataset  $D$ )
    - Put the example  $x$  into the training set:  $D_{train} = D_{train} \cup x$
    - Repeat the above 2 steps for  $n$  times
  - To use the set  $D_{train}$  for training the system
  - To use those examples in  $D$  **but not in**  $D_{train}$  to create the test set:  $D_{test} = \{z \in D; z \notin D_{train}\}$

# Bootstrap sampling (2)

- Important notes:
  - The training set has size of  $n$ , and an example in  $D$  may **appear multi times** in  $D_{train}$
  - The test set has size  $< n$ , and an example in  $D$  can **appear maximum to 1 time** in  $D_{test}$
- Suitable if we have a (very) small dataset  $D$

# Validation set

- The examples in the test set must not be used (in any way!) in the training of the system
- In some ML/DM problems, the system's training process includes 2 tasks:
  - Task 1: To train the system (= To learn approximately the target function)
  - Task 2: To optimize the values of the system's parameters
- The test set cannot be used for the purpose of optimization of the system's parameters
- To divide the whole dataset  $D$  into 3 disjoint subsets: *training set*, *validation set*, and *test set*
- The validation set is used to optimize the values of the system's parameters and the used ML/DM algorithm's ones
  - For a parameter, **the optimal value** is the one that results in **the best performance for the validation set**

# Evaluation metrics (1)

## ■ Accuracy

→ The accuracy degree of the prediction of the trained system to the test examples

## ■ Efficiency

→ The costs in time and memory resources needed for the training and the test of the system

## ■ Robustness

→ The tolerance degree of the system to noise/error/missing-value examples



# Evaluation metrics(2)

## ■ Scalability

→ How the system's performance (e.g., training/prediction speed) varies to the size of the dataset

## ■ Interpretability

→ How the system's results and operation are easy to understand for users

## ■ Complexity

→ The complexity of the model (i.e., the target function) learned by the system

# Select a trained model

- The selection of a trained model should compromise (balance) between:
  - The complexity of the trained model
  - The prediction accuracy degree of the trained model
- *Occam's razor*. A good trained model is the one that **is simple** and **achieves high accuracy (in prediction)** for the used dataset
- For example:
  - A trained classifier  $Sys1$ : (Very) simple, and rather (to a certain degree) fit to the training set
  - A trained classifier  $Sys2$ : More complex, and perfectly fit to the training set

→  $Sys1$  is preferred to  $Sys2$